



# INFS 519 – Fall 2015

## Program Design and Data Structures

### Lecture 8

Instructor: James Pope

As Lectured by: Arda Gumusalan

Email: [jpope8@gmu.edu](mailto:jpope8@gmu.edu)

# Today

- Questions from last class?
  - Binary Search Trees
  - Self Balancing Trees
  - AVL Rotations
- Schedule
  - 2-3 Trees
  - Red-Black Trees
  - B Trees

# Binary Search Tree

- A type of \_\_\_\_\_
- Relationship maintained between...
  -
- Rules
  - 
  - 
  - 
  -

# Binary Search Tree

- A type of binary tree!
- **Relationship** maintained between...
  - parent and both children
- **Rules**
  - parent  $>$  elements in left sub tree
  - parent  $<$  elements in right sub tree
  - both children are binary search trees
  - no duplicates (how do we handle this?)
- <http://people.ksp.sk/~kuko/bak/>

# Binary Search Tree: Big-O


- So... binary search trees
  - What is the **height**?
  - What is **worst/best/average** of:
    - **finding** an element
    - **inserting** an element
    - **deleting** an element
      - What if we want to delete the root?
      - Recall:
        - Predecessor: the largest value that is smaller than X.
        - Successor: the smallest values that is greater than X.

# How do we improve performance?

- What is the thing we need?
  - What causes the gap between worst case and average case?
- How do we get it?


# How do we improve performance?

- What do we need? Balance!
  - preferably **self-balancing!** (balance as you **add/remove/search** the tree)
- How? **Rotate!**




# Are These AVL Trees?


1




2




3




4




5




6



7




8






# Answers


1 Not AVL




2 AVL




3 Not AVL




4 AVL




5 Not AVL




6 AVL




7 Not AVL




8 AVL



# How do we fix this?




That's better!




# How do we fix this?

Will single rotation fix the following cases?

- An insertion in the right subtree of the left child of **X**
- An insertion in the left subtree of the right child of **X**




# Single Rotation Won't Fix!



(a) Before rotation

(b) After rotation

# Double Rotation



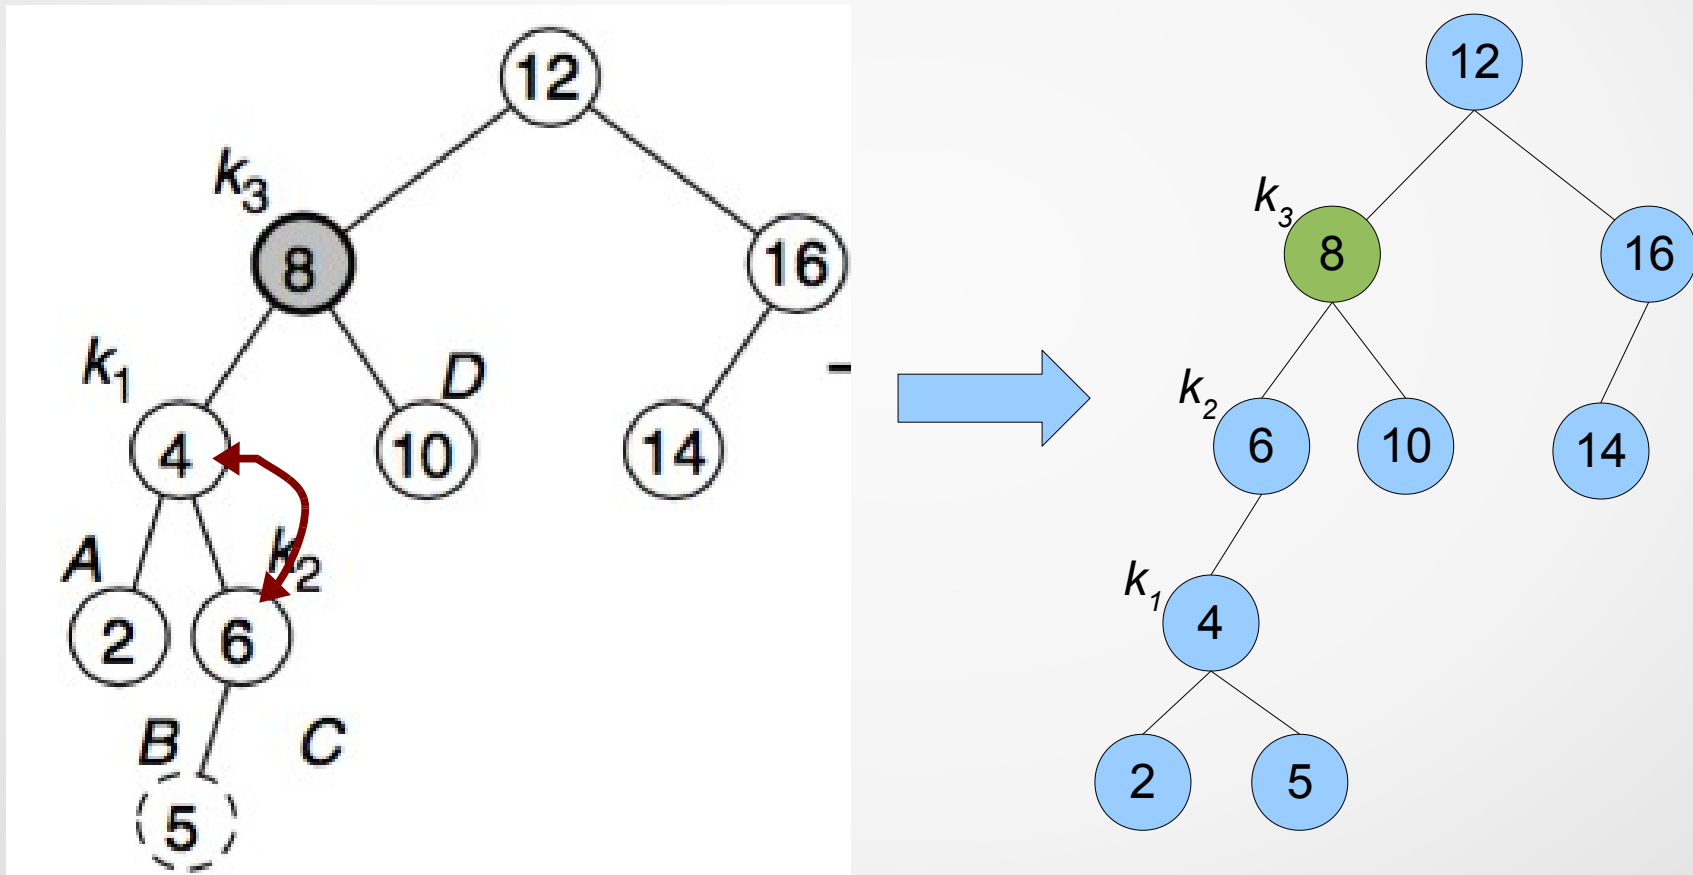
(a) Before rotation

(b) After rotation

- A rotation between  $X$ 's child and grandchild.
- A rotation between  $X$  and its new child.

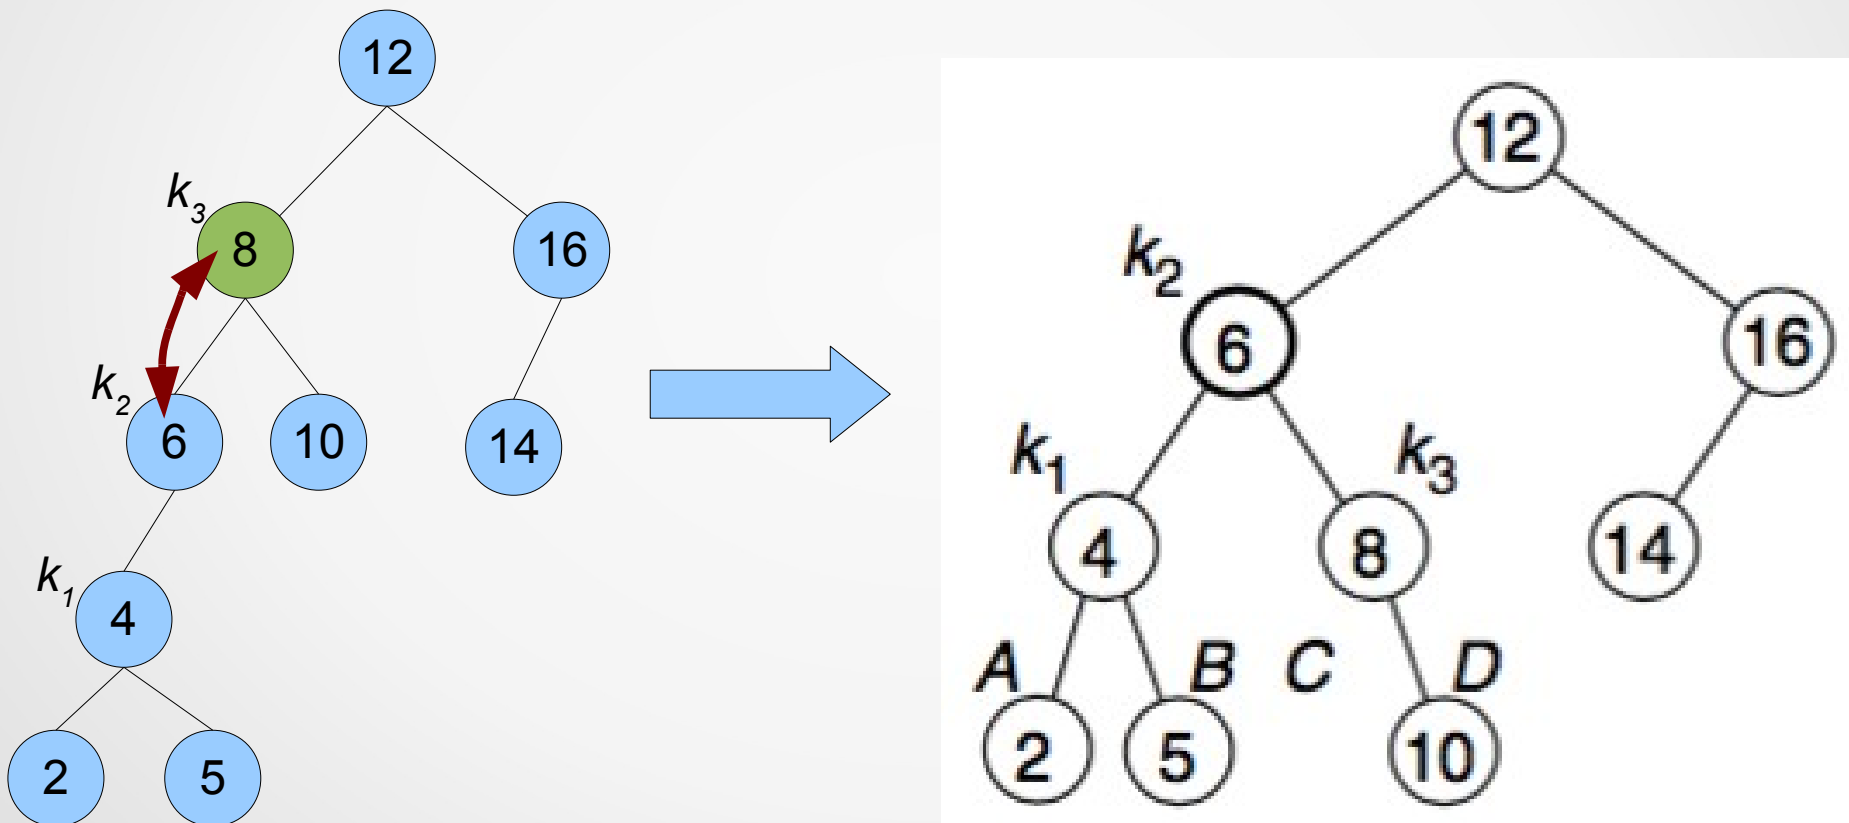
# Double Rotation-First Rotation

- A rotation between X's child and grandchild.



# Double Rotation-Second Rotation

- A rotation between X and its new child.







Questions?

# 2-3 Trees

Sedgewick/Wayne 3.3


- Invented by John Hopcroft
- In general, nodes are described by the **number of children** (i.e. number of links)
  - a **2-node** has 2 children
  - a **3-node** has 3 children
- Every node of BST is a 2-node
  - two links and one key

# 2-3 Tree Properties


Sedgewick/Wayne 3.3

- Every node of 2-3 tree is a 2 or 3 node
  - every **node** has between 1 and 2 **keys**
  - values are stored in **sorted order**
  - between 2 and 3 **children**
  - including **null links** (leaves only)
- 2-3 Trees are **perfectly balanced** – all null links (those of the leaves) are equal distance to the root

# 2-3 Trees: Order Property



- 2-node
  - 1 value (a), 2 subtrees (p, q)
  - $p < a < q$




- 3-node
  - 2 values (a, b), 3 subtrees (p, q, r)
  - $p < a < q < b < r$

images: <http://en.wikipedia.org/wiki/2%E2%80%93tree>

# 2-3 Tree Insert (Put) at Root

Sedgewick/Wayne 3.3

- Case1: Insert into a 2-node (no parent)
  - Simply add key to make it a 3-node




- Big clue: if your search ends at a 2-node, you always make it 3-node.

# 2-3 Tree Insert (Put) at Root

Sedgewick/Wayne 3.3




- Case2: Insert into a 3-node (no parent)
  - \_ Temporarily add the key (in order) to make a 4-node
  - \_ Take middle value, create the higher key node
  - \_ Create two new nodes, one with the left key and one with the right key
  - \_ Point the left child of the higher node to left key node and right child to right key node



# 2-3 Tree Insert (Put) with Parent

Sedgewick/Wayne 3.3


- Case3 Insert into a 3-node with 2-node parent
  - Similar to Case2, push middle key to parent



# 2-3 Tree Insert (Put) with Parent

Sedgewick/Wayne 3.3

- Case4 Insert into a 3-node with 3-node parent
  - Temporarily create 4-node, split as in Case2
  - Push middle up to parent, parent now 4-node
  - Push middle of parent up, split (harder)
  - Repeat (recursion to root if necessary)









# 2-3 Trees: Case 4

Sedgewick/Wayne 3.3

- Case4:



# 2-3 Tree Proposition

Sedgewick/Wayne 3.3

- Proposition: Insert and search operations into a 2-3 tree with  $N$  keys is guaranteed to visit at most  $\log(N)$  nodes.
- Proof: Consider two extremes for  $N$  keys (other scenarios fall between), when all nodes are 2-nodes and when all nodes are 3-nodes. Height of a 2-3 tree is between:
  - Lower bound (3-nodes):  $\text{floor}(\log_3(N))$
  - Upper bound (2-nodes):  $\text{floor}(\log_2(N))$

# Demo

- Gnarley trees.

# Red-Black Tree (RBT) Motivation

- Story so far
  - BST cannot guarantee performance for symbol table operations, specifically put, get, delete
  - AVL is balanced and relatively easy with casual implementation, but not as efficient as a 2-3 tree
  - 2-3 tree is (perfectly) balanced and can guarantee performance, more efficient
- So problem solved?
  - Recall one other desirable attribute for an algorithm, it should be easy to implement

# Red-Black Tree (RBT) Intuition

- BST easy to implement, 2-3 tree is balanced
  - Can we combine to get best of both?
  - “Yes”, store 2-3 tree in a BST structure
- Keep 1-to-1 correspondence between the implemented BST and the logically represented 2-3 tree
  - Put operation involves several different cases
  - Remove/delete also maintain invariant
- Debatable whether RBT is easy, though it is easier than 2-3 tree and performs well

# Red-Black Tree as a 2-3-4 Tree

Weiss 19.5

- 2-3-4 trees are also balanced
- Weiss presents RBT using 2-3-4 approach
- Invariants using 2-3-4 approach
  - 1) Every node colored either red or black
  - 2) Root is black
  - 3) If node is red, its children must be black
  - 4) Every path from a node to a null link must contain the same number of black nodes
- Involves several rotation cases

# Red-Black Tree as a 2-3 Tree


Sedgewick 3.3

- Sedgewick presents RBT using 2-3 approach
- Invariants using 2-3 approach
  - 1) Red links lean left (“left leaning RBT”)
  - 2) Root is black
  - 3) No node has two red links connected to it
  - 4) Tree has perfect black balance: every path from the root to a null link has the same number of black links
- Also involves several rotation cases. We will use this approach for RBT.


# Red-Black as 2-3 Tree and BST

Sedgewick/Wayne 3.3


red-black BST



horizontal red links



2-3 tree





1-1 correspondence between red-black BSTs and 2-3 trees



# RBT: Maintaining Search Order and Perfect Black Balance

- Insert nodes at bottom with red link
- Left rotations and right rotation operations
- Flip color operation
  - Flip children (i.e. red to black)
  - Flip parent from black to red




# RBT: Insert into 2-node

- Insert into a single 2-node  $\rightarrow$  root


## Left insert

left



## Right insert



right



# RBT: Insert into 2-node

- Insert into 2-node at bottom.



## Right insert







- Left insert does not require a rotation.

# RBT: Insert into a single 3-node





larger



smaller



between

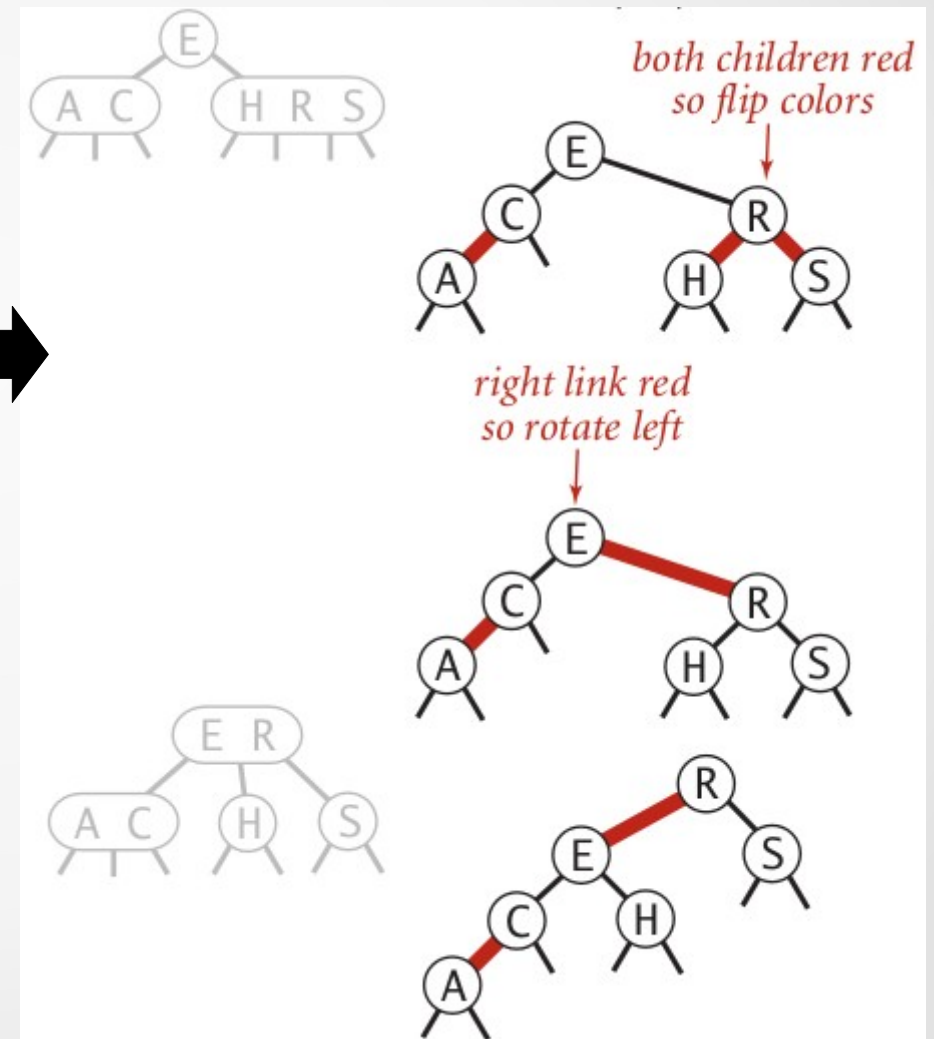



# RBT: Root Color

- Root should always be black. If red, means it is part of 3 node (i.e. not the root).
- Color flips can change the root's color to be red. In this exceptional case, the root should be changed to black.
  - When the root is changed from red to black, it means that the black height has increased by one

# RBT: Insert into 3-node at bottom

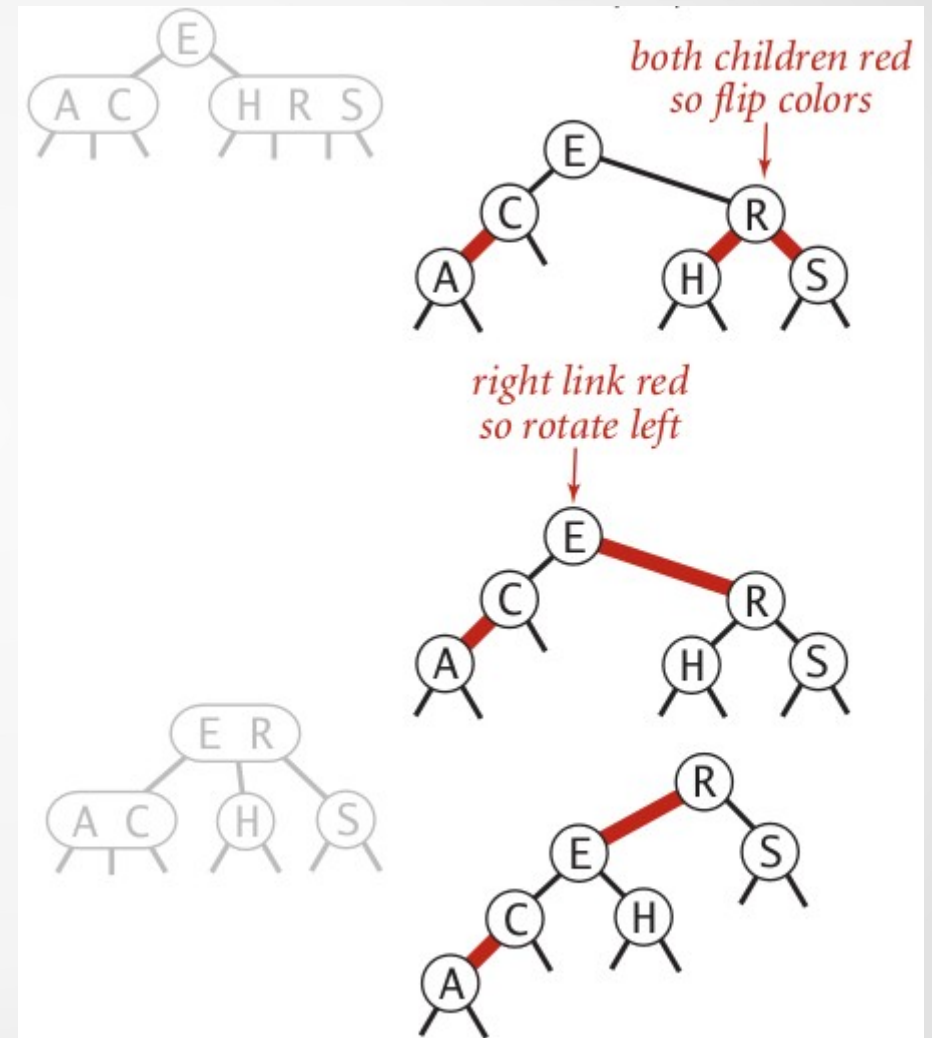
inserting H






# RBT: Passing Red Link Up Tree

- Rotation transformations are local but may result in a red link being passed up to its parent
- Can treat recursively
  - Think of it as adding a new node/red link.
- Repeat until parent is a 2 node or the root



# Left Leaning RBT: Summary

- Can maintain RBT as 2-3/BST **recursively** by the following rules:
  - 1) If right child is red and left child is black, then rotate left
  - 2) If left child is red and its left child is red, then rotate right
  - 3) If both left and right child are red, flip colors
- RBT a 2-3 Tree and BST. The get operation is same as BST.
  - \_ One of the main reasons for all this trouble...





# Symbol Table Summary

- Generally use hash table unless guaranteed performance or need ordered operations

Implementation	Worse-Case		Average-Case		Order Ops	remarks
	Search	Insert	Search	Insert		
Unordered List	N	N	N	N	No	
Ordered Array	$\lg N$	N	$\lg N$	N	Yes	
BST	N	N	$\lg N$	$\lg N$	Yes	Easy
AVL	$\lg N$	$\lg N$	$\lg N$	$\lg N$	Yes	Easy
<b>Red-Black</b>	<b><math>\lg N</math></b>	<b><math>\lg N</math></b>	<b><math>\lg N</math></b>	<b><math>\lg N</math></b>	<b>Yes</b>	<b>Often Used</b>
HT Chaining	N	N	N / M	N / M	No	Often Used
HT Probing	N	N	1	1	No	

\* Depending on variant, will assume  $O(\lg(N)) \sim O(1)$



Questions?

# Huge Data Sets

- Up to now entire data structure fits in memory
  - If it's too big it won't **fit into memory**...
  - ... so **use a tree** to break it up, store on disk
- Now, we have to perform in-memory instructions intermixed with disk accesses
  - Can ~25 million instructions in one second
  - Can ~ 6 disk accesses in one second
    - Creates a big bottleneck.
- Given 1 million records, assuming disk access required, balanced BST
  - Requires 20 accesses, about 3.5 seconds

# Partition Keys / Values

- Only need to keep keys (or subset of keys) in memory
- Data stored on disk and accessed when needed
  - Accessing disk performs best when reading/writing disk **blocks** (sometimes referred to as **pages**) at a time.

# B Tree Motivation

- Problem of data too big for memory was solved by breaking data into tree and storing much of the data on disk.
  - Created new problem. Now have to access disk
  - Asymptotically nothing has changed but the constants, even for  $\lg(N)$  operations is objectionable
- How can we reduce the constant (said another way, how can we reduce the height of the tree)?

# B Tree Motivation

- Organize data within the tree to correspond to pages on disk.
  - One tree node represents one disk page.
  - Read one disk page at a time.
  - Maximum degree chosen depends on disk characteristics and problem description.
- If we read 1000 keys, but look at 10, is that good?
  - Yes, the biggest latency is before we get the first bit, consecutive reads are relatively faster.
  - We end up making only  $1/10$  as many disk reads.

# B Tree versus B+ Trees

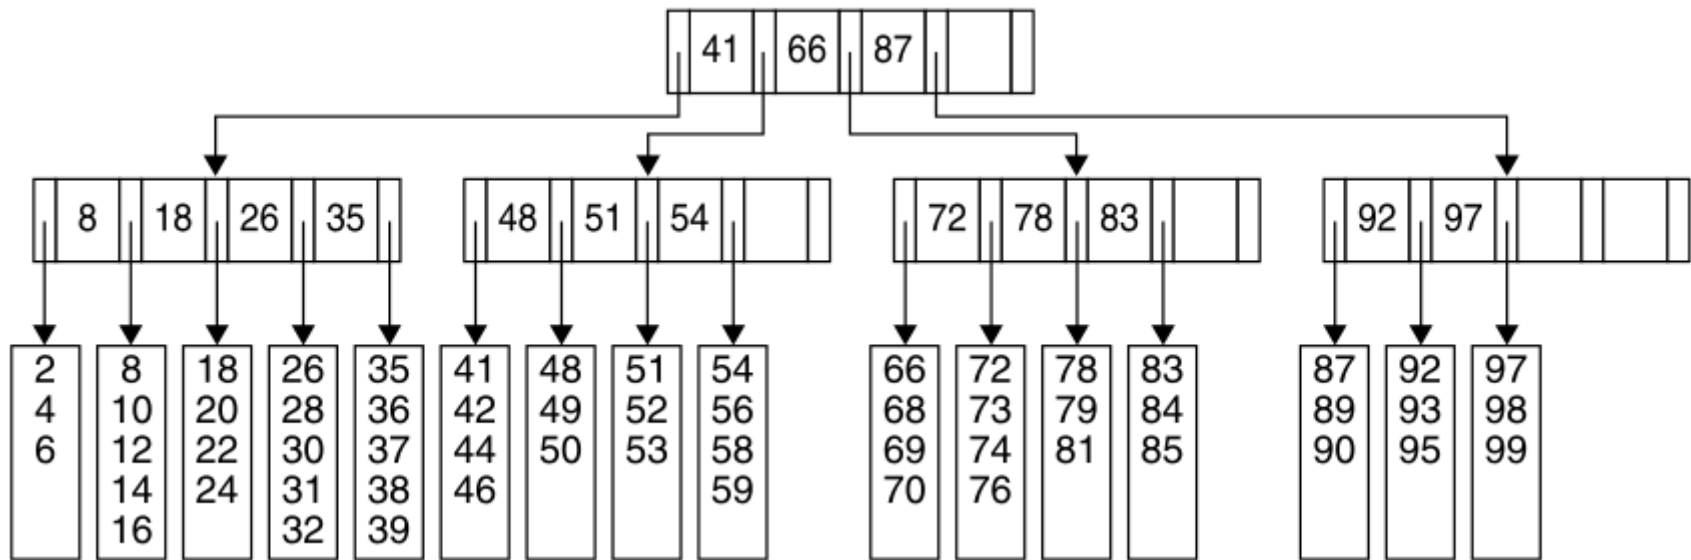
- B Trees
  - internal nodes store keys and values
  - Think of it as 2-3-....-many nodes
- B+ same as B Trees, except
  - internal nodes do not store values
    - only store keys, i.e. stores index like book
  - sibling leaves may have a pointer to link them in order
- Weiss 19.8 presents “B Tree” but is technically a B+ Tree. Will use term “B Tree” to refer to “B+ Tree” and only distinguish if necessary.

# B Trees

- Commonly taught, commonly used, easy to implement
- Basic idea:
  - Tree + List = B-Tree
  - We want a really big list...
    - but if it's too big it won't fit into memory...
    - ... so use a tree to break it up
- When to fix balance?
  - inserting/deleting



## 2 for 1: It's a list! In a tree!




# Learning B Trees

- 2-3 Trees
  - Restricted, simpler version of B Trees
- B Trees
  - General form of 2-3 trees
- B+ Trees
  - Same idea, different internal nodes

# B Tree Parameters

- Usually referred to by the M value
  - so an order 5 B-tree, has  $M=5$
- Additional Note: a 2-3 tree is a B Tree with:
  - $M = 3$
  - $L = 2 \rightarrow$  number of data items

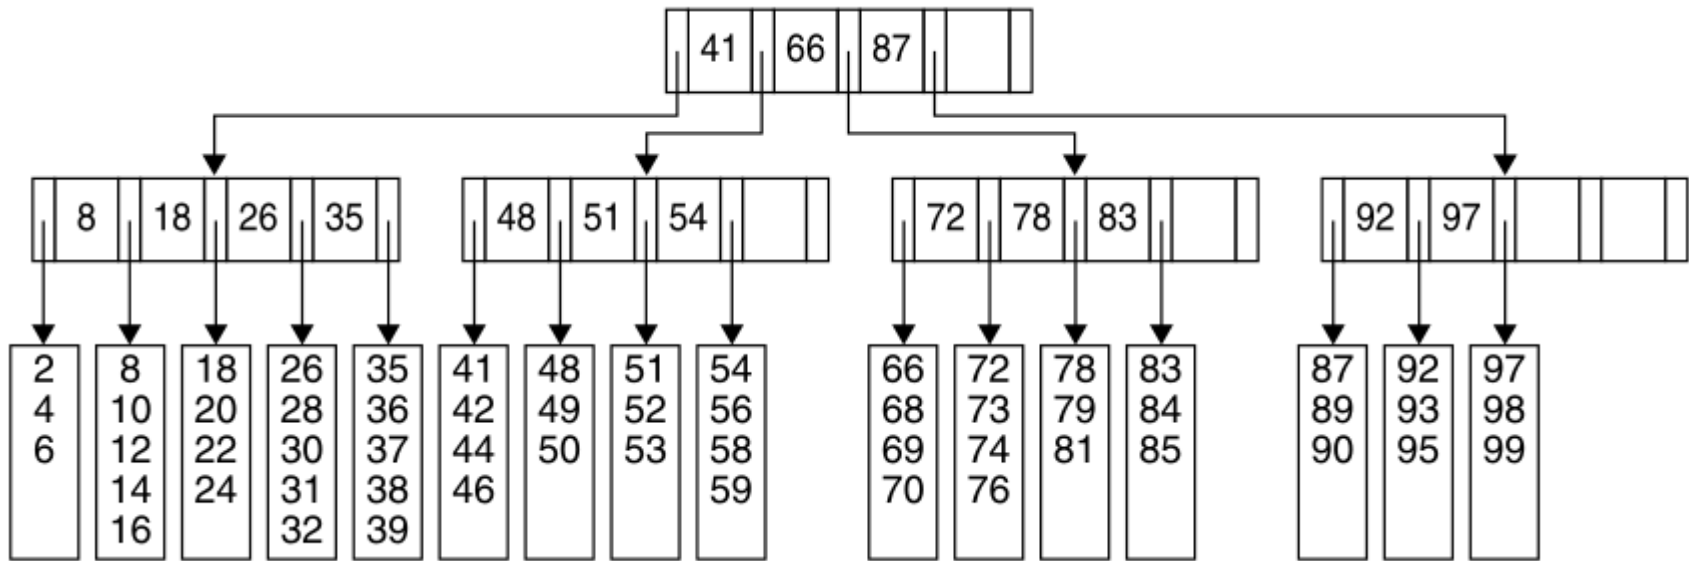


$M=5$

# B Tree Properties

- B tree of order M is an M-ary tree, invariants:
  - 1) Data Items are stored as leaves
  - 2) Non-leaf node store M-1 keys to guide search
  - 3) Root is leaf or between 2 and M children
  - 4) Non-leaf nodes (other than root) have between  $\text{ceil}(M/2)$  and M children
  - 5) All leaves are at same depth and have  $\text{ceil}(L/2)$  and L data items
- Note: Invariants keep tree from becoming degenerate

# B Tree Properties

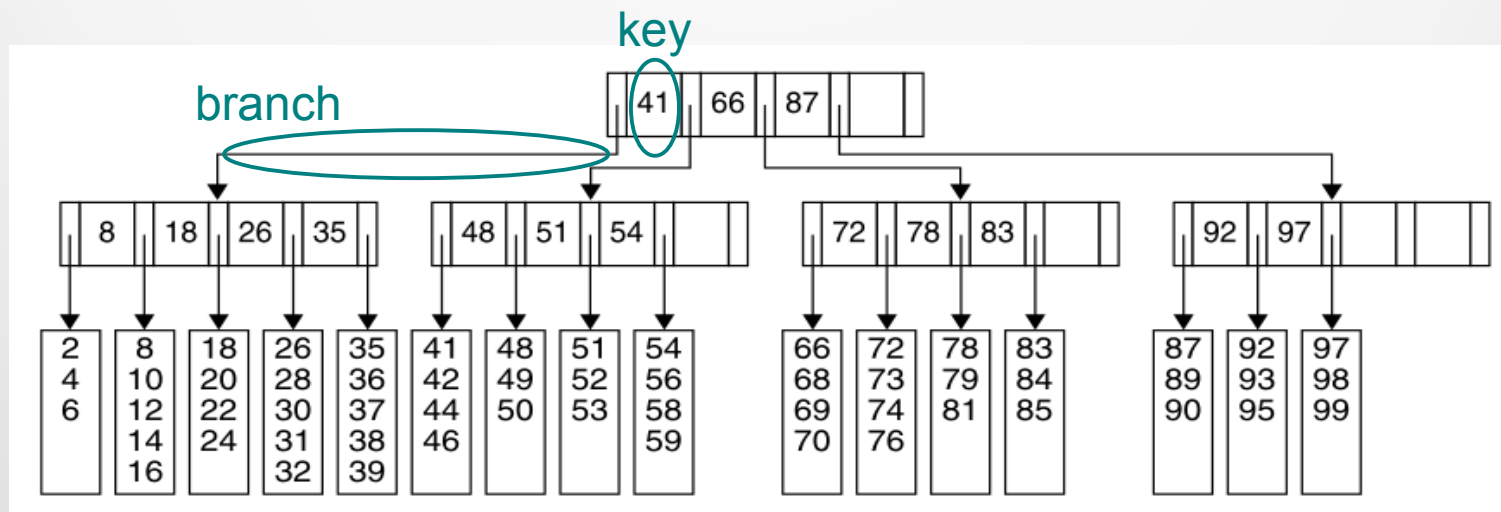


- A tree of order 5
- All nonleaf nodes have between 3 and 5 children
  - Which also means between 2 and 4 keys
- Each leaf node has between 3 and 5 data items
  - $\text{ceil}(L/2) \leq \# \text{ of data items} \leq L$

# B Tree Example Calculations 1/2

Weiss, Page 758

- Choose the maximum M and L that allow internal and external nodes to fit in one disk block. Example, store 500,000,000 records:
  - Assume block size is 4096 bytes, each key uses 8 bytes, each branch uses 4 bytes, each data record uses 32 bytes.



# B Tree Example Calculations 1/2

Weiss, Page 758

- Block size=4096 bytes, Key=8 bytes, Branch=4 bytes, Data record=32 bytes
- $M=?$ 
  - $M$  is for non-leaf nodes
    - $M-1$  keys  $\rightarrow 8(M-1)$
    - $M$  branches  $\rightarrow 4M$
    - Adds up to  $12M-8$  bytes
      - $12M-8=4096 \rightarrow M=342$
  - $L$  is for leaf nodes
    - Will store only data records.
      - $4096/32 \rightarrow L=128$

# B Tree Example Calculations 2/2

Weiss, Page 758

- Internal nodes branch at least  $\text{ceil}(342/2)=172$
- $\text{Log}_{172}(500 \text{ million}) = 3.89$ , max tree height 4
- At most  $500 \text{ million} / 64 = 7,812,500$  leaves
- In general:
  - Worse case number of accesses  $\sim \log_{M/2}(N)$
  - At most leaves is  $= N / \text{ceil}(L / 2)$



# Profound Implication

[https://en.wikipedia.org/wiki/Observable\\_universe](https://en.wikipedia.org/wiki/Observable_universe)

- The number of atoms on the observable universe
  - Approximately  $10^{80}$  (roughly  $2^{266}$ )
- B Tree with  $M=2048$  nodes
  - Max height  $\sim \log_{1024}(N) = \log_{1024}(2^{266})$   
 $\log_2(2^{266}) / \log_2(1024) \sim 26.6$
- Assuming you had the disk storage, did you anticipate that in two dozen steps you could search for any atom in the universe?

# B Tree vs other Balanced Trees

- B Tree reduces height of tree and therefore potentially number of disk accesses
- However, if all data can fit **in-memory**, then other balanced trees (e.g. RBT) should be used
  - Generally constants to traverse/insert into a RBT are better than B Tree



Questions?

# PA8

- Extend your binary search tree into a balanced AVL tree



Free Question Time!